

# Programming with SQL

Bringing Data to Your Codebase

# Section I Overview

# Agenda

---

Here is what we will do for the next 3 hours:

- 1 Expectations and Setup
- 2 Reading Data in Python, R, and Java
- 3 Writing Data in Python, R, and Java
- 4 Connection Management and Design Strategy

# What to Expect

---

**We are going to learn how to leverage SQL from programming platforms like Python, R, and Java.**

- While we will not cover each functionality encyclopedically, we will learn enough functionalities to create fully functional database applications.
- We will also cover connection pooling, design strategy, and other considerations when making a SQL database talk to a coded application.



# What Not to Expect

---

**We obviously cannot cover programming with SQL on every platform like Go, Swift, Rust, .NET, Julia, or *<put favorite language here>***

- Hopefully the knowledge we gain from using Python, R, and Java will give a good starting point to transfer knowledge to these other platforms.
- The principles and design philosophy we learn here should be largely translatable from these three platforms.



# Setting Up Python

---

**If you want to follow along, have your favorite Python environment set up:**

<https://www.python.org/>

**Code files and other resources are here:**

[https://github.com/thomasnield/oreilly\\_programming\\_with\\_sql](https://github.com/thomasnield/oreilly_programming_with_sql)

Make sure to have the *thunderbird\_manufacturing.db* database file in the working folder.

**Have the following libraries set up:**

- SQLite (already included with Python library)
- SQLAlchemy
- Pandas



# Setting Up R

---

**Have your favorite R environment set up:**

<https://www.r-project.org/>

<https://rstudio.com/>

**Code files and other resources are here:**

[https://github.com/thomasniel/oreilly\\_programming\\_with\\_sql](https://github.com/thomasniel/oreilly_programming_with_sql)

Make sure to have the *thunderbird\_manufacturing.db* database file in the working folder.

**Have the following packages set up:**

- DBI
- RSQLite
- pool

# Setting Up Java

---

**Have your favorite Java 8 (or later) environment set up:**

<https://aws.amazon.com/corretto/>

<https://www.jetbrains.com/idea/>

**Code files and other resources are here:**

[https://github.com/thomasniel/oreilly\\_programming\\_with\\_sql](https://github.com/thomasniel/oreilly_programming_with_sql)

Make sure to have the *thunderbird\_manufacturing.db* database file in the working folder.

**Have the following libraries set up (Maven or Gradle build system recommended):**

- com.zaxxer:HikariCP:3.4.1
- org.xerial:sqlite-jdbc:3.30
- tech.tablesaw:table-saw-core:0.36.0





# Section II

## Reading Data in Python, R, and Java

# EXERCISE

---

**Use Python, R, or Java to retrieve PRODUCT records with a PRICE of at least 100 and print them.**

## EXERCISE – Python Solution

---

```
from sqlalchemy import create_engine, text

engine = create_engine('sqlite:///thunderbird_manufacturing.db')
conn = engine.connect()

stmt = text("SELECT * FROM PRODUCT WHERE PRICE >= 100")
results = conn.execute(stmt)

for record in results:
    print(record)
```

## EXERCISE – R Solution

---

```
library(DBI)
library(RSQLite)

db <- dbConnect(SQLite(), dbname='thunderbird_manufacturing.db')

my_query <- dbSendQuery(db, "SELECT * FROM PRODUCT WHERE PRICE >= 100")

my_data <- dbFetch(my_query, n = -1)

dbClearResult(my_query)

print(my_data)

remove(my_query)
dbDisconnect(db)
```

# EXERCISE – Java Solution

---

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.ResultSet;
import java.sql.Statement;

public class JavaLauncher {

    public static void main(String[] args) {

        try {
            Connection conn = DriverManager.getConnection("jdbc:sqlite:/c:/my_folder/thunderbird_manufacturing.db");
            Statement stmt = conn.createStatement();
            ResultSet rs = stmt.executeQuery("SELECT * from PRODUCT WHERE PRICE >= 100");

            while (rs.next()) {
                System.out.println(rs.getString("PRODUCT_NAME") + " " + rs.getString("PRICE"));
            }
            //release connection
            conn.close();

        } catch (Exception e) {
            e.printStackTrace();
        }

    }
}
```

# Section III

## Writing Data in Python, R, and Java

# EXERCISE

---

**Use Python, R, or Java to create a new PRODUCT with a name of “NiteHawk”, a product group of “BETA”, and a price of “41”.**

**However, do it in a transaction but never complete the transaction. Do a query of the PRODUCT table to see if it was inserted temporarily.**

# EXERCISE – Python Solution

---

```
from sqlalchemy import create_engine, text

engine = create_engine('sqlite:///thunderbird_manufacturing.db')
conn = engine.connect()

# Create a transaction, but do not commit it
transaction = conn.begin()

# INSERT a new record
stmt = text("INSERT INTO PRODUCT (PRODUCT_NAME,PRODUCT_GROUP,PRICE) VALUES (:productName,:productGroup,:price)")

conn.execute(stmt, productName="NiteHawk",
              productGroup="BETA",
              price=41.0
             )

# Check records to see if last one inserted
for r in conn.execute(text("SELECT * FROM PRODUCT")):
    print(r)

# Close connection, don't commit
conn.close()
```



## EXERCISE – R Solution

---

```
library(DBI)
library(RSQLite)

db <- dbConnect(SQLite(), dbname='thunderbird_manufacturing.db')

my_query <- dbSendQuery(db, "SELECT * FROM PRODUCT WHERE PRICE >= 100")

my_data <- dbFetch(my_query, n = -1)

dbClearResult(my_query)

print(my_data)

remove(my_query)
dbDisconnect(db)
```

# EXERCISE – Java Solution

---

```
import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.ResultSet;
import java.sql.Statement;

public class JavaLauncher {

    public static void main(String[] args) {

        try {
            Connection conn = DriverManager.getConnection("jdbc:sqlite:/c:/my_folder/thunderbird_manufacturing.db");
            Statement stmt = conn.createStatement();
            ResultSet rs = stmt.executeQuery("SELECT * from PRODUCT WHERE PRICE >= 100");

            while (rs.next()) {
                System.out.println(rs.getString("PRODUCT_NAME") + " " + rs.getString("PRICE"));
            }
            //release connection
            conn.close();

        } catch (Exception e) {
            e.printStackTrace();
        }

    }
}
```

# Section IV

## Pooling and Design Strategy

# Design Strategy - Onus of Work

---


**When using SQL with a programming platform like Python, Java, or R, you will constantly be making a decision where the onus of processing will happen.**

**Should the database engine do the computation work, or the programming platform?**

- You can simply pull in data and have your Python/Java/R codebase do the heavy-lifting.
- You can also leverage more complex SQL against the database, and have Python/Java/R consume the results.
- With a very large, expensive and calculated dataset you can save it to a temporary table and use it to support your Python/R/Java application.

**A good rule of thumb: start with the simplest solution with minimal code/SQL that liberally hits the database as-needed, and gradually introduce caching strategies as performance starts to warrant it.**

**Never concatenate parameters, and use established SQL libraries to inject parameters safely to prevent SQL injection.**



# Design Strategy – Connection Pooling

---

**Modern applications (whether they are client side or server side), can be enormously complex.**

- Code design and modules change.
- Several services can be running simultaneously.

**Typically multiple threads are executing code simultaneously to parallelize work.**

- Having multiple threads accessing a single database connection can be hazardous, unless only one thread accesses and uses the connection at a time.
- It can be helpful to have multiple connections to support multiple threads safely.
- It is more efficient to reuse a connection rather than create/dispose it for each task.



# Design Strategy – Connection Pooling

---

Because most applications are **multithreaded**, or running several tasks simultaneously, it can be necessary to use a **connection pool** that maintains several threads.

- Connection pools are set up to maintain a fixed number of connections, and reuse them as needed.
- Of course, threads and database connections are expensive, so we must be conservative in how many we allow.
- If there are more threads than there are connections, then threads may be queued to wait for an available database connection.
- If a connection expires, the thread pool will dispose and replace it.

**For example, we can choose to have up to five database connections in a connection pool, and have threads use them as needed.**

**Connection pools can be set up dynamically to increase and decrease the number of connections based on needs.**



# Connection Pooling Libraries

---

**Java** – HikariCP (fastest and recommended), Vibur, TomCat, C3PO

**Python**- SQLAlchemy

**R** – Pool

# Design Strategy – Error Handling

---

One thing we did not talk about is **error handling**, and how this can cause connection leaks.

When a block of code fails using a database connection, it may skip over the line of code releasing the database connection.

To handle this, we can always make sure to dispose the connection in a **try-catch-finally** block so the connection is let go regardless if the operation is successful.

When available, libraries and specialized language features can automatically handle the disposal of connections in a block of code as well (e.g. Java has try-with-resources).





# Preventing SQL Injection

---

To prevent SQL injection, *never* concatenate a SQL string with parameters

Instead, use the right tools and libraries to safely inject parameters for you

*For Python, use SQLAlchemy*

```
from sqlalchemy import create_engine, text


engine = create_engine('sqlite:///C:\\Users\\thoma\\Dropbox\\rexon_metals.db')
conn = engine.connect()

def customer_for_id(customer_id):
    stmt = text("SELECT * FROM CUSTOMER WHERE CUSTOMER_ID = :id")
    return conn.execute(stmt, id=customer_id).fetchone()

print(customer_for_id(2))
```

*More info at:*

<http://www.sqlalchemy.org/>



# Preventing SQL Injection

---

*For Java, Scala, Kotlin, and other JVM languages use JDBC's PreparedStatement*

```
int customerId = 2;

Connection connection =
    DriverManager.getConnection("jdbc:sqlite:C:\\Users\\thoma\\Dropbox\\rexon_metals.db");

String sql = "SELECT * FROM CUSTOMER WHERE CUSTOMER_ID = ?";

PreparedStatement ps = connection.prepareStatement(sql);
ps.setInt(1, customerId);

ResultSet rs = ps.executeQuery();
rs.next();

System.out.println(rs.getInt("CUSTOMER_ID") + " " + rs.getString("NAME"));


connection.close();
```

*More info at:*

<http://tutorials.jenkov.com/jdbc/index.html>

<http://www.marcobehler.com/make-it-so-java-db-connections-and-transactions>

# SQL Injection Humor



Source: Google Images

Source: <https://xkcd.com/327/>

# SQL Injection in the News

---

## Simple Voice-Command SQL Injection Hack into Alexa Application

<https://securityboulevard.com/2019/09/simple-voice-command-sql-injection-hack-into-alexa-application/>

## How a 'NULL' License Plate Landed One Hacker in Ticket Hell

<https://www.wired.com/story/null-license-plate-landed-one-hacker-ticket-hell/>

## This couple cannot do the simplest things online because their last name is 'Null'

<https://thenextweb.com/insider/2016/03/27/last-name-null-is-tough-for-computers/>

# EXERCISE

---

**If you intend on having thousands of simultaneous tasks being executed in an application, it is a good idea to create a connection for each task (TRUE/FALSE)**

# EXERCISE

---

**If you intend on having thousands of simultaneous tasks being executed in an application, it is a good idea to create a connection for each task (TRUE/FALSE)**

FALSE! Creating thousands of connections to process thousands of tasks can be computationally expensive and cause crashes on the application side and server side. It is better to pool a limited number of connections (e.g. 12-24 connections) and reuse those connections to process the tasks.

# EXERCISE

---

**Which of the following are benefits of connection pools?**

- A) They reuse connections making the application more efficient
- B) Timed out connections will automatically get disposed and replaced
- C) Thread safety is enforced with database connection resources
- D) They will automatically take threads back from tasks when tasks are complete.
- E) Multiple database queries/updates can be executed simultaneously and safely

# EXERCISE

---

**Which of the following are benefits of connection pools?**

- ✓ A) They reuse connections making the application more efficient
- ✓ B) Timed out connections will automatically get disposed and replaced
- ✓ C) Thread safety is enforced with database connection resources
- ✗ D) They will automatically take threads back from tasks when tasks are complete.
- ✓ E) Multiple database queries/updates can be executed simultaneously and safely

**All the above are benefits of connection pools except for item “D”. The coded task is responsible for releasing a connection back to the connection pool so another task can use it.**



# HOMEWORK

---

**Use your knowledge of SQL, SQLite, and your favorite programming language to build a personal banking app.**

- A) Be able to enter bank transactions (with a date, description, and currency amount) and save them to a SQLite database.
- B) Create a table of categories (grocery, utilities, mortgage, etc) and be able to attach them to the transactions
- C) Create spending reports by category across day/week/month/year.

**Depending on your comfort, you can choose to make this application completely in a command-line environment, a desktop application, or an HTML frontend.**

## Other Online Trainings by Thomas Nield

---

[\*SQL Fundamentals for Data\*](#)

[\*Intermediate SQL for Data Analytics\*](#)

[\*Intro to Mathematical Optimization\*](#)

[\*Machine Learning from Scratch\*](#)